

MICROSERVICES LAB

EXERCISES ON KUBERNETES with solution

Create a simple deployment of the given app with name of your choice and 3 replicas of pods. Check the status of pod by sending request. App should be accessed from outside the cluster. Use Port Forwarding to Access Applications in a Cluster.

dep.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: usn-nginx-deployment
  labels:
    app: usn-nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: usn-nginx
  template:
    metadata:
      labels:
        app: usn-nginx
    spec:
      containers:
        - name: nginx
          image: 172.1.14.168:5001/nginx
          ports:
            - containerPort: 80
```

Command to deploy:

```
kubectl apply -f dep.yaml
```

Command to check pods

```
kubectl get pods -l=app=usn-nginx
```

Command to expose

```
kubectl expose deployment usn-nginx-deployment --type=NodePort --name=usn-nginx-service
```

To get exposed port

```
kubectl get svc -l=app=usn-nginx
http://172.1.14.168:<nodeport>
```

```
kubectl port-forward deployment/<dep_name>/<svc_name> newport:oldport
```

```
kubectl port-forward deployment/usn-nginx-deployment/usn-nginx newport:<nodeport>
```

Demonstrate the updation of image in live container in a pod using command line as well as by updating yaml files

dep.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: usn-nginx-deployment
  labels:
    app: usn-nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: usn-nginx
  template:
    metadata:
      labels:
        app: usn-nginx
    spec:
      containers:
        - name: nginx
          image: 172.1.14.168:5001/nginx
          imagePullPolicy: "Always"
          ports:
            - containerPort: 80
```

Command to deploy:

```
kubectl apply -f dep.yaml
```

Command to expose

```
kubectl expose deployment usn-nginx-deployment --type=NodePort --name=usn-nginx-service
```

Command to update image:

```
kubectl set image deployment/usn-nginx-deployment nginx=newImage
```

To check the updated name:

```
kubectl describe deploy usn-nginx-deployment | grep Image:
```

Create busybox pod with two containers, each one will have the image busybox and will run the 'sleep 3600' command. Make both containers mount an emptyDir at '/etc/foo'. Connect to the second busybox, write the first column of '/etc/passwd' file to '/etc/foo/passwd'. Connect to the first busybox and write '/etc/foo/passwd' file to standard output. Delete pod.

<pre> dep.yaml apiVersion: v1 kind: Pod metadata: creationTimestamp: null labels: run: busybox name: 172.1.14.168:5001/busybox spec: dnsPolicy: ClusterFirst restartPolicy: Never containers: - args: - /bin/sh - -c - sleep 3600 image: busybox imagePullPolicy: IfNotPresent name: 172.1.14.168:5001/busybox resources: {} volumeMounts: - name: myvolume mountPath: /etc/foo - args: - /bin/sh - -c - sleep 3600 image: busybox name: busybox2 volumeMounts: - name: myvolume mountPath: /etc/foo volumes: - name: myvolume emptyDir: {} </pre>	<p>Command to deploy: kubectl apply -f dep.yaml</p> <p>Connect to the second container:</p> <pre> kubectl exec -it busybox -c busybox2 -- /bin/sh cat /etc/passwd cut -f 1 -d ':' > /etc/foo/passwd cat /etc/foo/passwd exit </pre> <p>Connect to the first container:</p> <pre> kubectl exec -it busybox -c busybox -- /bin/sh mount grep foo # confirm the mounting cat /etc/foo/passwd exit </pre> <p>To delete pod</p> <pre> kubectl delete po busybox </pre>
-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

Perform the following.

- a. Create 3 pods with names nginx1, nginx2,nginx3. All of them should have the label app=v1 Show all labels of the pods.
- b. Get only the 'app=v2' pods.
- c. Remove the 'app' label from the pods we created before

```
kubectl run usn-nginx1 --image=nginx --restart=Never --labels=app=usn-v1
```

```
kubectl run usn-nginx2 --image=nginx --restart=Never --labels=app=usn-v1
```

```
kubectl run usn-nginx3 --image=nginx --restart=Never --labels=app=usn-v1
```

```
kubectl get po --show-labels
```

```
kubectl get po -l app=usn-v2
```

```
kubectl label po nginx1 nginx2 nginx3 app-
```

Create a deployment with image nginx:1.7.8, called nginx, having 2 replicas, defining port 80 as the port that this container exposes

1. Check how the deployment rollout is going
2. Update the nginx image to nginx:1.7.9
3. Check the rollout history and confirm that the replicas are OK
4. Undo the latest rollout and verify that new pods have the old image (nginx:1.7.8)
5. Do an on purpose update of the deployment with a wrong image nginx:1.91
6. Verify that something's wrong with the rollout
7. Return the deployment to the second revision (number 2) and verify the image is nginx:1.7.9
8. Check the details of the fourth revision

<pre>dep.yaml apiVersion: apps/v1 kind: Deployment metadata: name: usn-nginx-deployment labels: app: usn-nginx spec: replicas: 2 selector: matchLabels: app: usn-nginx template: metadata: labels: app: usn-nginx spec: containers: - name: nginx image: 172.1.14.168:5001/nginx ports: - containerPort: 80</pre>	<pre>kubectl apply -f dep.yaml kubectl set image deploy usn-nginx-deployment nginx=nginx:1.7.9 kubectl rollout history deploy usn-nginx-deployment kubectl rollout undo deploy usn-nginx-deployment # wait a bit # select one 'Running' Pod kubectl get po usn-nginx-deployment kubectl describe po pod-name grep -i image # should be nginx:1.7.8 kubectl set image deploy usn-nginx-deployment nginx=nginx:1.91 kubectl rollout status deploy usn-nginx-deployment kubectl rollout undo deploy usn-nginx-deployment --to-revision=2 kubectl describe deploy usn-nginx-deployment grep Image: kubectl rollout status deploy usn-nginx-deployment # Everything should be OK</pre>
----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

How to expose multiple port in kubernetes services or Multi-Port Services

deploy_ports.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: multiport-demo
  namespace: default
spec:
  replicas: 2
  selector:
    matchLabels:
      app: multiport-demo
  template:
    metadata:
      labels:
        app: multiport-demo
    spec:
      containers:
        - name: multiport-demo
          image: 172.1.14.168:5001/nginx
```

```
apiVersion: v1
kind: Service
metadata:
  name: multiport-demo
  namespace: default
spec:
  type: NodePort
  selector:
    app: multiport-demo
  ports:
    - port: 80
      name: http
      targetPort: 80
      nodePort: 30030
    - port: 443
      name: https
      targetPort: 80
      nodePort: 30031
```

```
kubect1 apply -f deploy_ports.yaml
http://172.1.14.168:30030/
http://172.1.14.168:30031/
```

From the above example the multiport-demo service will be exposed internally to cluster applications on port 80 and externally to the cluster on the node IP address on 30030,30031. It will also forward requests to pods with the label “app: multiport-demo” on port 80 from 80,443

Example of Multi-Container Pod

deploy_multicontainer.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: mc1
spec:
  volumes:
  - name: html
    emptyDir: {}
  containers:
  - name: 1st
    image: 172.1.14.168:5001/nginx
    volumeMounts:
    - name: html
      mountPath:
/usr/share/nginx/html
  - name: 2nd
    image:
172.1.14.168:5001/debian
    volumeMounts:
    - name: html
      mountPath: /html
    command: ["/bin/sh", "-c"]
    args:
    - while true; do
      date >> /html/index.html;
      sleep 1;
    done
```

```
kubectl apply -f
deploy_multicontainer.yaml
kubectl describe po mc1
kubectl exec mc1 -c 1st --
/bin/cat
/usr/share/nginx/html/index.
html
kubectl exec mc1 -c 2nd --
/bin/cat /html/index.html
kubectl exec --stdin --tty
mc1 -c hello -- /bin/sh
```

Create a Pod with ubuntu image and a command to echo "YOUR_NAME" which overrides the default CMD/ENTRYPOINT of the image.

dep_ubuntu_pod1.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu
  labels:
    app: ubuntu
spec:
  containers:
  - name: ubuntu
    image: 172.1.14.168:5001/ubuntu
    command: ["/bin/bash"]
    args: ["-c", "echo MSRIT"]
```

```
kubectl apply -f dep_ubuntu_pod1.yaml
kubectl logs ubuntu
kubectl exec --stdin --tty ubuntu -- /bin/bash
kubectl delete pod ubuntu
```

Create a Pod that runs one container. The configuration file for the Pod defines a command and arguments by using environment variables:

dep_ubuntu_pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu
  labels:
    app: ubuntu
spec:
  containers:
  - name: ubuntu
    image: 172.1.14.168:5001/ubuntu
    env:
      - name: MESSAGE
        value: "MSRIT"
    command: ["/bin/echo"]
    args: ["$(MESSAGE)"]
```

```
kubectl apply -f dep_ubuntu_pod.yaml
kubectl logs ubuntu
kubectl delete pod ubuntu
```

Create a Pod that runs two Containers. The two containers share a Volume that they can use to communicate.

dep_2container_pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: two-containers
spec:
  restartPolicy: Never
  volumes:
  - name: shared-data
    emptyDir: {}
  containers:
  - name: nginx-container
    image: 172.1.14.168:5001/nginx
    volumeMounts:
    - name: shared-data
      mountPath: /usr/share/nginx/html
  - name: debian-container
    image: 172.1.14.168:5001/debian
    volumeMounts:
    - name: shared-data
      mountPath: /pod-data
    command: ["/bin/sh"]
    args: ["-c", "echo Hello from the debian container > /pod-data/index.txt"]
```

```
kubectl apply -f dep_2container_pod.yaml
kubectl get pods
kubectl exec -it two-containers -c nginx-container -- /bin/bash
curl localhost
```